

# Assignment 5: Beer Sales

*Joshua Goldberg*

*May, 09 2019*

Load data from TSA package (the package is written by authors Jonathan Cryer and Kung-Sik Chan).

```
data(beersales)
```

The data is the monthly beer sales in millions of barrels, 01/1975 - 12/1990.

Train: 01/1975 - 12/1989.

Test: 1990

## Part 1

Use ARIMA(p,d,q) model to forecast beer sales for all months of 1990 using the following two multi-step forecasting approaches. For each model, check mean, autocorrelation and normality of the residuals. Confirm if the residuals are white noise.

```
train <- window(beersales, c(1975, 1), c(1989, 12))
test <- window(beersales, c(1990, 1), c(1990, 12))
```


```
arima_fit <- auto.arima(train)
summary(arima_fit)
```

```
## Series: train
## ARIMA(4,1,2)(2,1,2)[12]
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ma1      ma2      sar1      sar2
##          0.5103 -0.1662  0.1032 -0.3966 -1.1757  0.3125  0.6838 -0.592
## s.e.      0.1453  0.0986  0.0863  0.0789  0.1493  0.1421  0.1451  0.165
##          sma1      sma2
##          -1.1967  0.5849
## s.e.      0.1394  0.2087
##
## sigma^2 estimated as 0.2837:  log likelihood=-134.55
## AIC=291.1   AICc=292.81   BIC=325.4
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.01226384 0.4974721 0.3570516 -0.1687359 2.535728 0.6855892
##              ACF1
## Training set 0.001560233
```

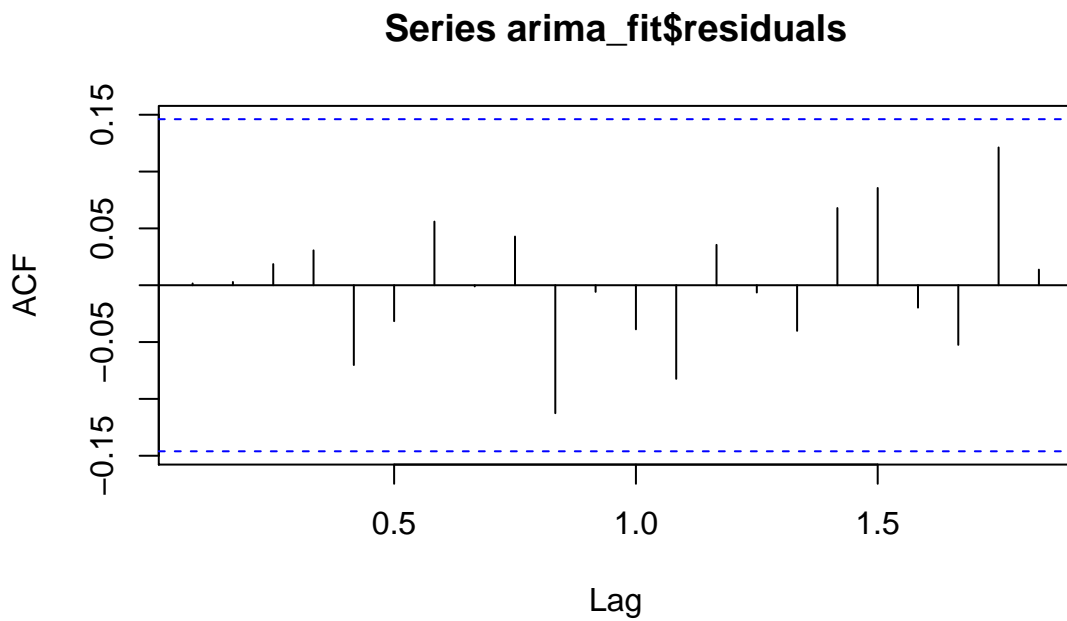
## 1A

Use the h-period in `forecast()` to forecast each month of 1990. This is also known as recursive forecasting where you fit a model only once and use it recursively for h-periods.

```
ts.plot(arima_fit$residuals)
```



```
acf(arima_fit$residuals)
```



```
(recursive_forecast <- forecast(arima_fit, h = 12))
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 1990	13.81601	13.13331	14.49871	12.77191	14.86011
## Feb 1990	13.07707	12.35715	13.79698	11.97605	14.17808
## Mar 1990	14.96181	14.23546	15.68817	13.85095	16.07268
## Apr 1990	15.58503	14.83785	16.33220	14.44232	16.72774
## May 1990	17.24847	16.49698	17.99996	16.09917	18.39777
## Jun 1990	16.86360	16.10993	17.61727	15.71096	18.01624
## Jul 1990	16.95571	16.19987	17.71156	15.79974	18.11168
## Aug 1990	17.02231	16.26451	17.78012	15.86336	18.18127
## Sep 1990	14.28619	13.51600	15.05638	13.10828	15.46410
## Oct 1990	14.55136	13.75967	15.34305	13.34057	15.76214
## Nov 1990	12.89695	12.09174	13.70216	11.66548	14.12841

```
## Dec 1990      12.30127 11.48554 13.11699 11.05372 13.54881
```

## 1B

Use the monthly data as a continuous time series. Forecast for 1990 Jan, Plug forecast into the time series, build a new model to forecast for 1990 Feb. And so on and so forth. In other words,  $h=1$  in all the forecasts. This is known as direct recursive (DirRec) forecasting where you fit a new model for each time step.

```
dir_rec <- function(.data, .model_count, refit = FALSE) {
  predictions <- vector("numeric", .model_count)
  models <- vector("list", .model_count)

  new_data <- train %>% as_tsibble(index = index) %>% append_row(.model_count)
  model_fit <- auto.arima(new_data %>% drop_na() %>% as_tsibble(index = index) %>% as.ts())
  orders <- arimaorder(model_fit)

  # Index used for appending new data
  index_change <- .model_count - 1

  for (i in 1:.model_count) {

    model_data <- new_data %>% drop_na() %>% as_tsibble(index = index) %>% as.ts()

    model <- Arima(model_data,
                   order = orders[c("p", "d", "q")],
                   seasonal = orders[c("P", "D", "Q")])

    forecast_point <- forecast(model, h = 1)$mean %>% as.numeric()

    predictions[[i]] <- forecast_point
    new_data[nrow(new_data) - index_change, "value"] <- forecast_point
    models[[i]] <- model

    # Reduce index by 1 floored to zero if negative
    index_change <- ifelse(index_change - 1 < 0, 0, index_change - 1)

    # Refit `auto.arima()` to obtain new parameters with prediction in consideration
    if (refit) {
      model_fit <- auto.arima(model_data, stepwise = FALSE)
      orders <- arimaorder(model_fit)
    }
  }

  list(data = new_data, models = models, predictions = predictions)
}


model_same_params <- dir_rec(train, 12, refit = FALSE)
model_refit <- dir_rec(train, 12, refit = TRUE)
```

## 1C


Plot the mean, the p-value of the autocorrelation test and the p-value of the normality test of the residuals of the 12 models. The Box test results fail to reject the null hypothesis (the data are independently distributed).

The data visually do not look too bad, but we reject the null hypothesis (data is normally distributed) of the shapiro test.

```
map_dbl(model_same_params$models, ~ Box.test(.x$residuals, lag = 24)$p.value) %>% plot()
```

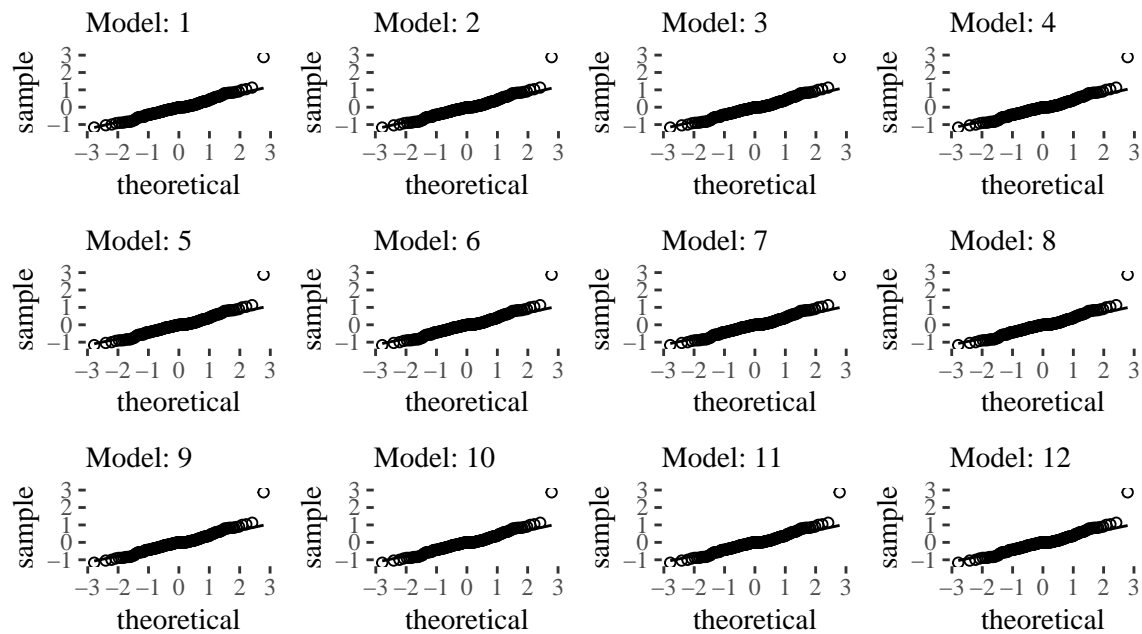


```
map_dbl(model_same_params$models, ~ shapiro.test(.x$residuals)$p.value) %>% plot()
```




```
map(1:12, ~
  data.frame(r = model_same_params$models[[.x]]$residuals %>% as.numeric()) %>%
  ggplot(aes(sample = r)) +
  stat_qq(pch = 1) +
  stat_qq_line() +
  labs(subtitle = paste0("Model: ", .x))) %>%
  patchwork::wrap_plots() +
  patchwork::plot_annotation(title = "Residual Diagnostics: QQ Plot")
```

## Residual Diagnostics: QQ Plot




```
walk(model_same_params$models, ~ checkresiduals(.x))
```

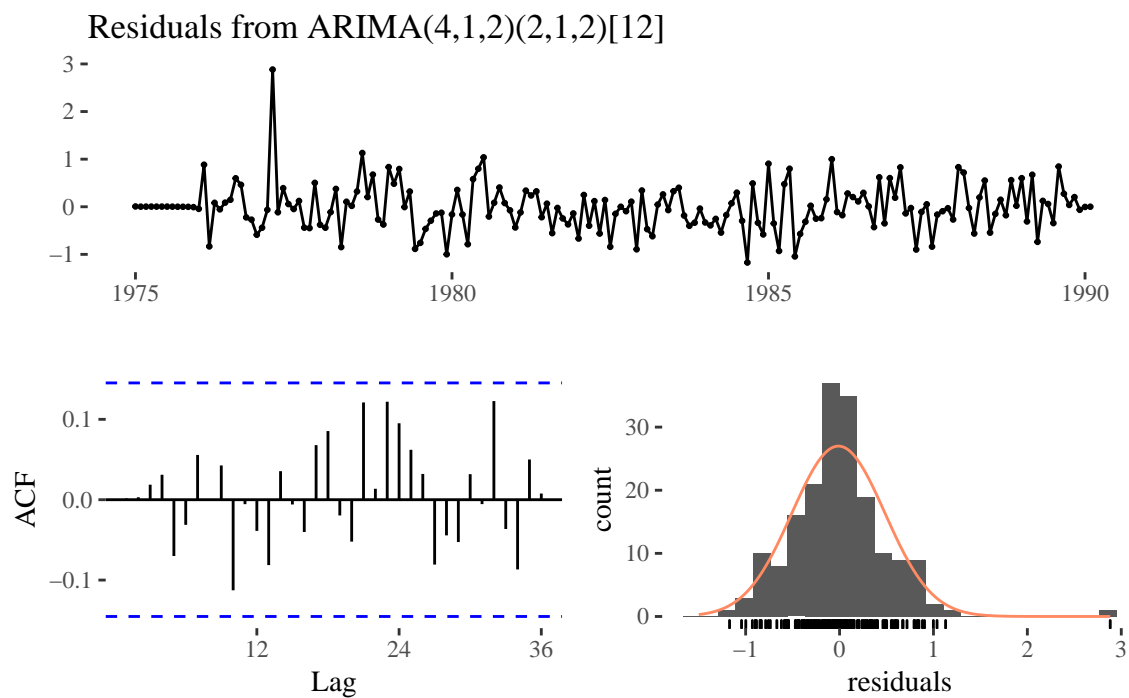
Residuals from ARIMA(4,1,2)(2,1,2)[12]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.035, df = 14, p-value = 0.2052
##
## Model df: 10.   Total lags used: 24
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.095, df = 14, p-value = 0.2025
##
## Model df: 10.   Total lags used: 24
```




```
##
##  Ljung-Box test
```


```
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.173, df = 14, p-value = 0.199
##
## Model df: 10. Total lags used: 24
```



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.253, df = 14, p-value = 0.1955
##
## Model df: 10. Total lags used: 24
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.315, df = 14, p-value = 0.1928
##
## Model df: 10.   Total lags used: 24
```




```
##
##  Ljung-Box test
```




```
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.406, df = 14, p-value = 0.1889
##
## Model df: 10. Total lags used: 24
```



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.478, df = 14, p-value = 0.1859
##
## Model df: 10. Total lags used: 24
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.535, df = 14, p-value = 0.1835
##
## Model df: 10.   Total lags used: 24
```




```
##
##  Ljung-Box test
```


```
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.618, df = 14, p-value = 0.1801
##
## Model df: 10. Total lags used: 24
```



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.676, df = 14, p-value = 0.1777
##
## Model df: 10. Total lags used: 24
```



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.756, df = 14, p-value = 0.1745
##
## Model df: 10.   Total lags used: 24
```



```
##
##  Ljung-Box test
```


```
##
## data: Residuals from ARIMA(4,1,2)(2,1,2)[12]
## Q* = 18.831, df = 14, p-value = 0.1715
##
## Model df: 10. Total lags used: 24
```

## Part 2

Plot the Recursive and DirRec along with the actuals. Use `ylim=c(12.5, 17)` to get a good visual of the plot differences.

```
test_predictions <- test %>%
  as_tsibble() %>%
  mutate(direct_recursive = model_refit$predictions,
         recursive_model = recursive_forecast %>% as_tibble() %>% pull(`Point Forecast`))

test_predictions %>%
  rename(actual = value) %>%
  gather(key = type, value = value) %>%
  ggplot(aes(index, value, color = type)) +
  geom_point() +
  geom_line() +
  scale_color_viridis_d(name = NULL, labels = c("Actual", "Direct Recursive", "Recursive")) +
  labs(title = "Direct Recursive (refitting arima each run) vs. Recursive",
       x = "Date",
       y = "Sales")
```



## Part 3

Calculate the MSE for 1990 - which of the two approaches take larger computation time and why? Direct Recursive (DR) is better from . Computationally it is more intensive, since it has to refit the model with new data.

```
mse_recursive = mean((test_predictions$recursive_model - test_predictions$value)^2)
mse_direct_recursive = mean((test_predictions$direct_recursive - test_predictions$value)^2)

glue::glue("MSE recursive: {mse_recursive}")

## MSE recursive: 0.565002004429162

glue::glue("MSE direct recursive: {mse_direct_recursive}")

## MSE direct recursive: 0.352798601820935
```